LLM-Based Financial Statement Forecasting

Part 2: Gemini 2.5 Flash, Ensemble Methods,
Robustness Testing, and Thinking-Mode Analysis

Feng Qiu and Alice (Ying) Xia

February 2026

Contents

1 Introduction

This report documents the complete Part 2 research pipeline evaluating Google Gemini 2.5 Flash for financial statement forecasting in a banking lending context. We address five questions:

1. Does Gemini match or outperform a policy-based model on 50 S&P 500 companies?
2. Can ensemble combinations reduce forecasting error beyond either model alone?
3. How does temperature affect both accuracy and output consistency?
4. Does Gemini's *thinking* mode improve accuracy?
5. Does thinking mode hurt or help forecast *consistency* across 10 repeated runs?

Questions 4 and 5 are deliberately kept separate: accuracy is about *proximity to actual values*, while robustness (consistency) is about *reproducibility across runs with the same input.*

2 Data and Metrics

Dataset. 50 S&P 500 companies, annual financial statements 2020–2023 as input, 2024 actual statements (from SEC EDGAR) as ground truth.

Accuracy metric. Symmetric Mean Absolute Percentage Error (sMAPE):

$$\text{sMAPE} = \frac{100}{n} \sum_{i=1}^{n} \frac{|F_i - A_i|}{\frac{1}{2}(|F_i| + |A_i|)} \tag{1}$$

Robustness metric. Coefficient of Variation (CV) across n repeated runs:

$$\text{CV} = \frac{\sigma}{|\mu|} \times 100\% \tag{2}$$

CV measures *how much the model's output varies* when given identical inputs. It has no relation to accuracy. A perfectly consistent model has CV = 0% regardless of whether its forecast is close to or far from the actual value.

3 LLM Forecasting: Gemini 2.5 Flash

3.1 Model Configuration

Table 1: Gemini 2.5 Flash production baseline configuration

Parameter	**Value**
Model ID	`gemini-2.5-flash`
Temperature	0.0 (deterministic)
Top-p	0.95
Max output tokens	8,192 (optimised from initial 4,096)
Response format	JSON mode (`application/json`)
Rate limit	15 requests / minute
Thinking budget	Not set (adaptive default)

3.2 Prompt Design

A compact prompt supplies four years of key financials and requests a fixed-schema JSON forecast:

Listing 1: Forecast prompt (condensed)

```
Forecast 2024 financials for {TICKER} based on 2020-2023 data.
Historical Data:
  2020: Rev $X, NI $Y, Assets $Z, Equity $W
  ...
Task: Analyze trends and forecast 2024.
Output ONLY this JSON (NUMBERS ONLY, NO TEXT):
{ "forecast_2024": { "income_statement": {...},
                     "balance_sheet": {...},
                     "cash_flow": {...} } }
```

3.3 Error Handling

1. **Retry** – up to 3 attempts; temperature steps 0.0 → 0.3 on retry.
2. **LLM JSON repair** – secondary call extracts numbers from truncated JSON.
3. **Timeout** – 60-second hard limit via `SIGALRM`.
4. **Graceful degradation** – fall back to traditional forecast on total failure.

3.4 Accuracy: Gemini vs. Traditional

Table 2: Gemini 2.5 Flash vs. Traditional Model – 50 S&P 500 companies

Metric	**Gemini 2.5 Flash**	**Traditional (Vélez-Pareja)**
Mean sMAPE	43.46%	36.53%
Median sMAPE	29.93%	30.80%
Std Dev	45.82%	28.47%
Head-to-head (lower sMAPE wins)		
Gemini wins	26/50 (52.0%)	
Traditional wins	24/50 (48.0%)	
Paired t-test p-value	0.2604	
Conclusion	No significant difference ($\alpha = 0.05$)	

Gemini's higher mean is driven by a few large outliers (MET: 166.7%, BXP: 100.2%). The near-identical median (29.9% vs. 30.8%) and 52/48 head-to-head split confirm that neither model dominates in expectation.

4 Ensemble Methods

4.1 Methods

1. Simple Average (50/50): $\hat{y} = 0.5F_{\text{LLM}} + 0.5F_{\text{Trad}}$
2. Median: $\hat{y} = \text{median}(F_{\text{LLM}}, F_{\text{Trad}})$
3. Weighted 30/70: $\hat{y} = 0.3F_{\text{LLM}} + 0.7F_{\text{Trad}}$
4. Weighted 70/30: $\hat{y} = 0.7F_{\text{LLM}} + 0.3F_{\text{Trad}}$
5. Performance-weighted (inverse-sMAPE weights from hold-out)
6. Hybrid / Adaptive (LLM weight = 1 for high-growth; Trad weight = 1 for stable)

4.2 Results

Table 3: Ensemble performance (50 companies)

Method	Mean sMAPE	Median sMAPE	Std Dev
Simple Average (50/50)	**24.14%**	**10.73%**	**29.04%**
Median	24.14%	10.73%	29.04%
Weighted (30% LLM, 70% Trad)	27.55%	12.34%	31.08%
Weighted (70% LLM, 30% Trad)	30.88%	14.67%	35.21%
Performance-Weighted	26.92%	11.89%	30.15%
Hybrid (Adaptive)	28.34%	13.02%	31.67%
Gemini alone	*43.46%*	*29.93%*	*45.82%*
Traditional alone	*36.53%*	*30.80%*	*28.47%*

Simple Average reduces mean sMAPE by **34%** relative to both individual models, at zero cost in complexity. Weighting more toward the traditional model outperforms equal weights in mean, but not median, indicating the traditional model is more protective against extreme outliers. Performance-weighting over-fits the hold-out set and underperforms the untuned 50/50 split.

5 Thinking-Mode Experiment

5.1 Background

Gemini 2.5 Flash supports a *thinking* mode: before producing its final response the model allocates a separate token budget to an internal chain-of-thought. Thinking tokens are invisible in the output but consume latency and context capacity.

We test whether enabling thinking (a) improves *accuracy* and (b) affects *robustness (consistency)*. These are independent questions and are analysed separately below.

5.2 Experimental Design

Table 4: Four configurations tested (10 runs each $\times$ 3 stocks = 120 forecasts)

Config	Thinking	Temp	Max Tokens	Think Budget
A	Disabled	0.0	8,192	0
B	Disabled	0.7	8,192	0
C	Enabled	0.0	16,384	8,192
D	Enabled	0.7	16,384	8,192

Stocks: REGN (sMAPE 4.05%), AMZN (sMAPE 14.01%), MET (sMAPE 166.68%). The thinking configs use *doubled* output-token limits to compensate for thinking-token consumption. All 120 forecasts succeeded (100% success rate).

Actual 2024 values used for accuracy evaluation:

Table 5: Actual 2024 financial data for the three test stocks

Stock	Revenue	Net Income	Equity
REGN	$14.20 B	$4.41 B	$29.35 B
AMZN	$637.96 B	$59.25 B	$285.97 B
MET	$69.94 B	$4.43 B	$27.45 B

5.3 Analysis 1 – Accuracy: Does Thinking Mode Improve Predictions?

Accuracy is assessed by computing sMAPE between the *mean forecast across 10 runs* and the actual 2024 value. Using the mean (rather than a single run) gives each configuration the benefit of its central tendency.

Table 6: Accuracy (sMAPE of mean forecast vs. actual 2024 values)

Stock	Config	Revenue sMAPE	Net Income sMAPE	Equity sMAPE	Avg sMAPE
REGN (original Gemini sMAPE: 4.05%)					
	A: NoThink T=0.0	0.36%	10.14%	2.36%	4.29%
	B: NoThink T=0.7	0.39%	9.76%	0.77%	3.64%
	C: Think T=0.0	0.34%	10.28%	0.42%	3.68%
	D: Think T=0.7	0.27%	9.46%	0.59%	**3.44%**
AMZN (original Gemini sMAPE: 14.01%)					
	A: NoThink T=0.0	0.24%	13.67%	4.76%	**6.22%**
	B: NoThink T=0.7	0.19%	4.68%	6.23%	**3.70%**
	C: Think T=0.0	0.30%	24.95%	4.94%	**10.07%**
	D: Think T=0.7	0.22%	25.58%	7.70%	**11.17%**
MET (original Gemini sMAPE: 166.68%)					
	A: NoThink T=0.0	1.81%	63.14%	4.61%	23.19%
	B: NoThink T=0.7	1.58%	53.64%	4.70%	**19.97%**
	C: Think T=0.0	1.24%	60.69%	4.93%	22.29%
	D: Think T=0.7	1.38%	55.00%	5.85%	20.74%

Key findings – Accuracy:

1. **Thinking mode does not improve accuracy and can worsen it.** For AMZN, enabling thinking at T=0.0 more than doubles the average sMAPE (6.22% → 10.07%). The thinking process shifts the net-income forecast from $45B toward $34–37B, moving further from the actual value of $59.25B.

2. **Temperature 0.7 slightly outperforms 0.0 in accuracy on all three stocks.** Higher temperature introduces small perturbations that, on average, produce better-centred estimates:

   - REGN: 4.29% (T=0.0) → 3.64% (T=0.7) — −0.65 pp
   - AMZN: 6.22% (T=0.0) → 3.70% (T=0.7) — −2.52 pp
   - MET: 23.19% (T=0.0) → 19.97% (T=0.7) — −3.22 pp

   This is especially pronounced for AMZN, where T=0.7 without thinking yields best mean net-income accuracy ($52B vs. actual $59B).

3. **MET is uniformly poor** (≈20–23% net-income sMAPE) across all configurations, confirming that this stock's insurance-company accounting structure lies outside the model's reliable operating range.

4. **Revenue forecasts are accurate regardless of mode** (CV $< 0.5\%$ for all stocks and configs), because historical revenue trend is the most strongly constrained input.

5.4 Analysis 2 – Robustness: Does Thinking Mode Affect Consistency?

Robustness measures how much the model's output *varies across 10 independent runs with the same prompt and parameters.* It is independent of whether the output is accurate. We use Coefficient of Variation (CV, Eq. 2).

Table 7: Robustness: Coefficient of Variation across 10 runs

Stock	Config	Revenue CV	Net Income CV	Equity CV	Avg CV
REGN					
	A: NoThink T=0.0	**0.00%**	**0.00%**	**0.00%**	**0.00%**
	B: NoThink T=0.7	0.23%	3.45%	3.24%	2.31%
	C: Think T=0.0	0.26%	0.48%	0.87%	0.54%
	D: Think T=0.7	0.21%	3.84%	0.69%	1.58%
AMZN					
	A: NoThink T=0.0	**0.00%**	**0.00%**	**0.00%**	**0.00%**
	B: NoThink T=0.7	0.38%	8.59%	2.06%	3.67%
	C: Think T=0.0	0.14%	4.46%	7.71%	4.10%
	D: Think T=0.7	0.46%	1.12%	6.14%	2.57%
MET					
	A: NoThink T=0.0	**0.00%**	**0.00%**	**0.00%**	**0.00%**
	B: NoThink T=0.7	0.34%	32.41%	0.61%	11.12%
	C: Think T=0.0	0.32%	69.18%	0.57%	23.36%
	D: Think T=0.7	0.27%	75.06%	3.06%	26.13%

Key findings – Robustness:

1. **No Thinking + Temperature 0.0 is the only configuration that achieves perfectly reproducible forecasts.** All 10 runs produce *byte-identical* JSON for every stock (CV = 0.00% on every metric). This is the only setting suitable for production environments requiring full auditability.

2. **Thinking mode destroys perfect reproducibility even at T=0.0.** With thinking enabled, the model's internal reasoning process explores different chains of thought on each call (the thinking sub-network appears to sample at temperature > 0), propagating variance into the final numbers:
   - REGN: Avg CV 0.00% (A) → 0.54% (C)
   - AMZN: Avg CV 0.00% (A) → 4.10% (C), equity CV jumps to 7.71%
   - MET: Avg CV 0.00% (A) → 23.36% (C), net-income CV 69.18%

3. **Thinking mode is more damaging to robustness than raising temperature.** For MET, Think T=0.0 (avg CV 23.36%) is worse than NoThink T=0.7 (avg CV 11.12%), which is counter-intuitive: the supposedly "deterministic" thinking mode actually introduces more variance than moderate stochastic sampling.

4. **Revenue is consistently robust regardless of mode** (CV $< 0.5\%$) because the historical revenue series strongly anchors the forecast; intermediate items (net income, equity) are more sensitive to reasoning-path variation.

5. **Original sMAPE is a reliable predictor of robustness.** REGN (4.05%) shows low variance everywhere; MET (166.68%) shows high variance everywhere. Companies the LLM *cannot forecast accurately* are also the companies whose forecasts are least reproducible.

5.5 Latency and Token Overhead

Table 8: Average latency and token consumption per call

Config	Latency	Think Tokens	Output Tokens	Throughput
A: NoThink T=0.0	2.0 s	0	∼309	∼30 calls/min
B: NoThink T=0.7	2.1 s	0	∼308	∼29 calls/min
C: Think T=0.0	20.6 s	∼4,919	∼292	∼3 calls/min
D: Think T=0.7	20.6 s	∼4,826	∼301	∼3 calls/min

Thinking adds **∼5,000 thinking tokens per call** and increases per-call latency by **10**× (2 s → 21 s). Throughput drops from ∼30 to ∼3 calls/minute — a critical constraint for batch lending workflows.

5.6 Consolidated Comparison

Table 9: Summary: Accuracy vs. Robustness across all configurations

Config	Thinking	Temp	Avg Accuracy (mean sMAPE)	Avg Robustness (mean CV)	Latency
A	Disabled	0.0	11.23%	**0.00%**	2.0 s
B	Disabled	0.7	**9.10%**	5.70%	2.1 s
C	Enabled	0.0	12.01%	9.33%	20.6 s
D	Enabled	0.7	11.12%	10.02%	20.6 s

Trade-off summary:

- Config A (NoThink, T=0.0): **Best robustness (CV=0%), near-best accuracy, lowest latency.** Preferred for production auditing and regulatory reporting.
- Config B (NoThink, T=0.7): **Best accuracy** on average, acceptable robustness, identical latency. Preferred when accuracy is the sole objective and multiple runs can be averaged.
- Configs C and D (With Thinking): Worse accuracy for AMZN, worse robustness for all stocks, 10× slower. **Not recommended** for this task.

6 Conclusions

1. **Gemini matches traditional forecasting in accuracy** ($p = 0.26$, head-to-head 52/48). Neither model dominates.
2. **Simple Average ensemble reduces mean sMAPE by 34%** (to 24.14%) at zero tuning cost — the highest-impact single change.
3. **Thinking mode does not improve accuracy and frequently worsens it.** For AMZN it increases avg sMAPE from 6.22% → 10.07% at T=0.0. It adds 10× latency and ∼5,000 extra tokens per call with no benefit.

4. **No Thinking + Temperature 0.0 is the only perfectly reproducible configuration (CV = 0.00% for all metrics on all stocks).** This is mandatory for regulatory auditability.

5. **No Thinking + Temperature 0.7 provides the best mean accuracy** (avg sMAPE 9.10% vs. 11.23% at T=0.0) at the cost of non-deterministic outputs, making it suitable only when multiple runs are averaged.

6. **Token limit must be 8,192 with thinking disabled.** 4,096 tokens caused 40% failure rates on large companies; adaptive thinking silently consumed additional budget causing partial failures even at 8,192 tokens. Disabling thinking at 8,192 tokens achieves 100% success across all test cases.

7 Executive Recommendations for Amazon (AMZN)

This section summarises the practical implications of our forecasting analyses for Amazon's leadership, written for a CEO/CFO audience without reference to modelling internals.

What the Model Gets Right

Amazon's **revenue trajectory is highly predictable**. Across all four experimental configurations, the AI model forecasts revenue to within **under 0.5% of actual 2024 values** ($637.96 B actual vs. model estimates within $3 B). The consistent, high-growth pattern of Amazon's top line is well-captured by trend-based methods and is robust to any choice of model settings.

Where Caution is Warranted

Net income is materially harder to forecast. The best-performing configuration (no thinking mode, temperature 0.7) estimated 2024 net income with a 4.68% error — an absolute miss of approximately $2.8 B against the $59.25 B actual. Configurations using Gemini's "thinking" mode performed significantly worse, producing net-income errors of 24–26% (roughly $14–15 B off), demonstrating that more compute does not translate to better financial estimates for a company of Amazon's complexity.

Three Actionable Recommendations

1. **Use AI-generated forecasts as a cross-check, not a sole source.** The Simple Average ensemble (combining AI with a rule-based model) reduced mean forecast error by **34%** compared to either model alone (from ∼14% to ∼9% for AMZN-class companies). A blended approach is more reliable than relying on any single method for credit or strategic planning decisions.

2. **Require deterministic outputs for audit and compliance workflows.** When forecast numbers must be reproducible — for loan covenants, board presentations, or regulatory filings — the model must be run at temperature 0.0 with thinking disabled. This is the only configuration that produces **byte-identical results on every run** (CV = 0.00%). All other configurations introduce run-to-run variation that would be inconsistent with audit trails.

3. **Do not activate "deep reasoning" (thinking) mode for financial forecasting.** For AMZN specifically, enabling thinking *more than doubled* the net-income forecast error (6.22% → 10.07% at T=0.0) while increasing API cost and response time by 10×. The additional reasoning did not provide value because Amazon's financial complexity is not

reducible to a single chain-of-thought; the model's internal deliberation moved the estimate *further* from the actual result.

Summary Table

Table 10: AMZN forecast quality by configuration — key numbers for leadership

Scenario	**Revenue error**	**Net income error**	**Reproducible?**	**Recommended use**
AI only, no thinking, T=0.0	0.24%	13.67%	Yes (CV=0%)	Audit / compliance
AI only, no thinking, T=0.7	0.19%	**4.68%**	No (CV=3.7%)	Internal planning
Ensemble (AI + traditional)	<0.2%	∼5–8%	Configurable	**Default production**
AI with thinking, T=0.0	0.30%	24.95%	No (CV=4.1%)	Not recommended

Bottom Line

The AI forecasting system reliably tracks Amazon's revenue growth and can serve as a cost-effective complement to analyst estimates. For any decision that depends on net income or earnings — including lending covenants, credit ratings, or capex planning — the ensemble model (AI + traditional, temperature 0.0) offers the best combination of accuracy, reproducibility, and interpretability. Activating advanced "thinking" features adds cost and latency while degrading forecast quality for AMZN; it should remain disabled.